



Introduction to AI for postgraduate students

Lecture Note 3 Optimization


POSTECH



Reading Materials

- http://www.deeplearningbook.org/contents/linear_algebra.html
- <http://www.stat.cmu.edu/~ryantibs/convexopt/lectures/stochastic-gd.pdf>



Gradient Descent: Univariate



At a random point Start
Repeat

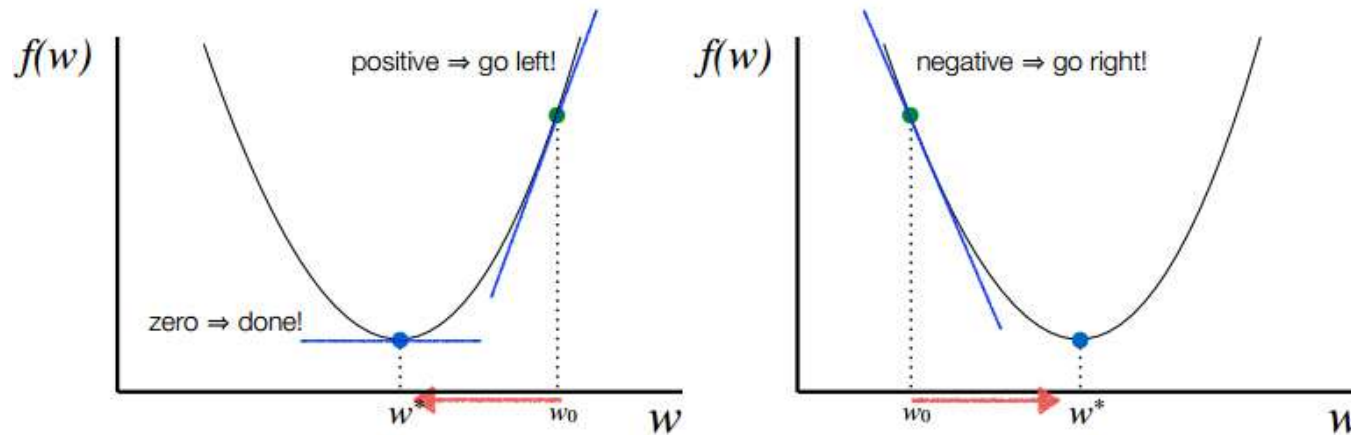
- Determine a descent direction
- Choose a step size Choose
- Update

Until some criterion is satisfied



Gradient Descent: Univariate

▪ Choosing the descent direction




▪ Update

Step Size

Update Rule: $w_{i+1} = w_i - \alpha_i \frac{df}{dw}(w_i)$

Negative Slope

Multivariate extension



Step Size


Update Rule: $w_{i+1} = w_i - \alpha_i \nabla f(w_i)$

Negative Slope


Gradient
e.g., $\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right)$

Gradient Descent: Univariate

- Choosing the step size





Too small: converge very slowly



Too big: overshoot and even diverge

Better way




Reduce size over time

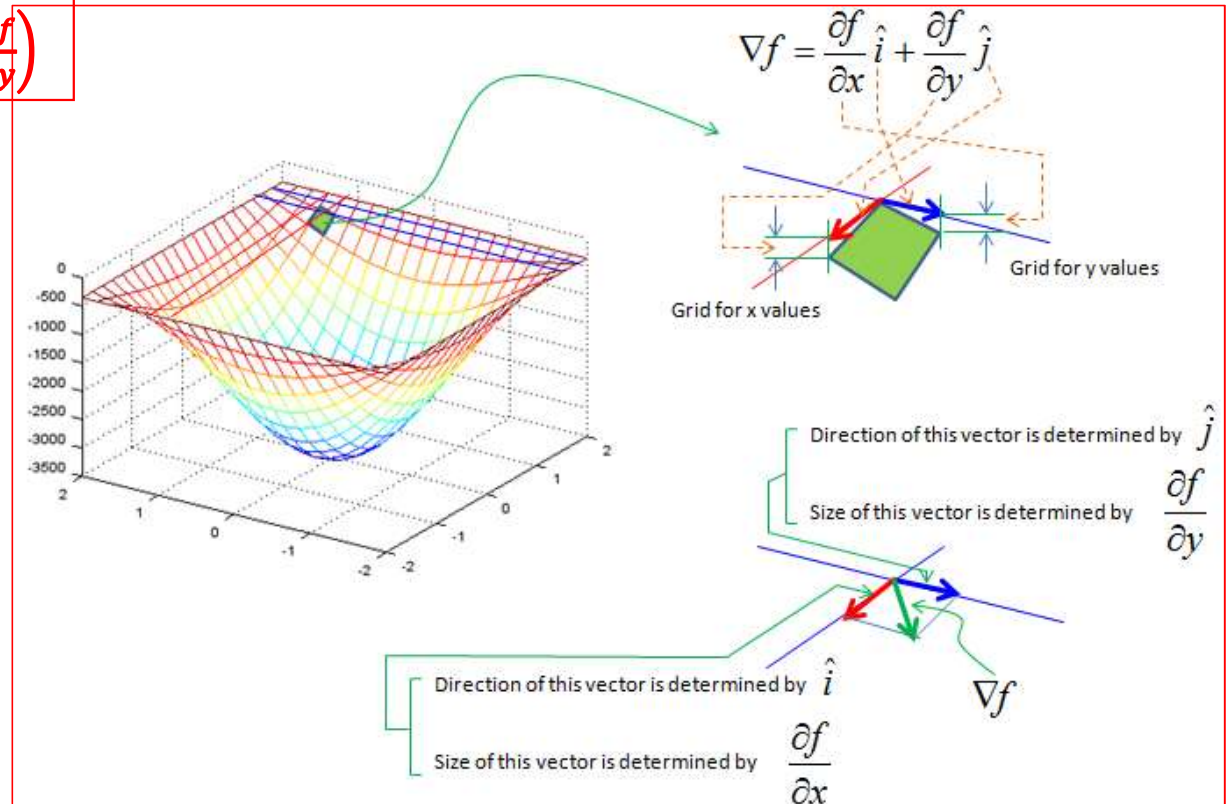

Gradient Descent: Multivariate

Definition of a Gradient vector

*Fig. from http://www.sharetechnote.com/html/Calculus_Gradient.html



$$\text{Or } \nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)$$



Gradient Descent: Multivariate

- Consider a 3-D function: $T(x, y, z)$
- Gradient vector is defined by: $\text{grad } T \equiv \left(\frac{\partial T}{\partial x}, \frac{\partial T}{\partial y}, \frac{\partial T}{\partial z} \right)$.
- By the chain rule:

$$dT = \frac{\partial T}{\partial x} dx + \frac{\partial T}{\partial y} dy + \frac{\partial T}{\partial z} dz.$$


$$d\mathbf{r} \equiv (dx, dy, dz)$$



$$dT = \text{grad } T \cdot d\mathbf{r}.$$


- Suppose that $dT=0$ for some $d\mathbf{r}$:

$$dT = \text{grad } T \cdot d\mathbf{r} = 0.$$





Gradient Descent: Multivariate

- Illustration of gradient vectors and contours



Limitations of Gradient Descent



Second Derivation Method: Gradient Descent

- **Directional derivative in direction “u” (unit vector)**

- Slope of the function f in direction \mathbf{u}

$$\frac{\partial}{\partial \alpha} f(\mathbf{x} + \alpha \mathbf{u})$$

- By the chain rule:

$$\frac{\partial}{\partial \alpha} f(\mathbf{x} + \alpha \mathbf{u}) \quad \longrightarrow \quad \mathbf{u}^\top \nabla_{\mathbf{x}} f(\mathbf{x}) \text{ when } \alpha = 0$$

- **We find the unit vector “u” such that the directional derivative is minimized → for steepest descent**

$$\min_{\mathbf{u}, \mathbf{u}^\top \mathbf{u} = 1} \mathbf{u}^\top \nabla_{\mathbf{x}} f(\mathbf{x})$$

$$= \min_{\mathbf{u}, \mathbf{u}^\top \mathbf{u} = 1} \|\mathbf{u}\|_2 \|\nabla_{\mathbf{x}} f(\mathbf{x})\|_2 \cos \theta$$

where θ is the angle between \mathbf{u} and the gradient

$$\longrightarrow \min_{\mathbf{u}} \cos \theta \longrightarrow$$

Solution: vector \mathbf{u} should be in the **opposite** direction of the Gradient

Second Derivation Method: Gradient Descent

- Steepest descent update

$$x' = x - \epsilon \nabla_x f(x)$$


Learning rate or
Step size

- Analytical discussion on choosing the step size
 - Read "4.3.1 Beyond the Gradient: Jacobian and Hessian Matrices"

Third Derivation Method: Gradient Descent

- Suppose that θ has two variables $\{\theta_1, \theta_2\}$

Given a point, we want to find the point with the smallest value nearby.



Third Derivation Method: Gradient Descent

- Multivariate Taylor series

$$\begin{aligned} h(x, y) = & h(x_0, y_0) + \frac{\partial h(x_0, y_0)}{\partial x} (x - x_0) + \frac{\partial h(x_0, y_0)}{\partial y} (y - y_0) \\ & + \text{something related to } (x - x_0)^2 \text{ and } (y - y_0)^2 \\ & + \dots \end{aligned}$$

When x and y are close to x_0 and y_0



$$h(x, y) \approx h(x_0, y_0) + \frac{\partial h(x_0, y_0)}{\partial x} (x - x_0) + \frac{\partial h(x_0, y_0)}{\partial y} (y - y_0)$$

Third Derivation Method: Gradient Descent

Based on Taylor Series:

If the red circle is *small enough*, in the red circle


$$L(\theta) \approx L(a, b) + \frac{\partial L(a, b)}{\partial \theta_1}(\theta_1 - a) + \frac{\partial L(a, b)}{\partial \theta_2}(\theta_2 - b)$$

$$s = L(a, b)$$

$$u = \frac{\partial L(a, b)}{\partial \theta_1}, v = \frac{\partial L(a, b)}{\partial \theta_2}$$

$$L(\theta)$$

$$\approx s + u(\theta_1 - a) + v(\theta_2 - b)$$



Third Derivation Method: Gradient Descent

Based on Taylor Series:

If the red circle is small enough, in the red circle

constant


$$s = L(a, b)$$

$$u = \frac{\partial L(a, b)}{\partial \theta_1}, v = \frac{\partial L(a, b)}{\partial \theta_2}$$

$$L(\theta) \approx s + u(\theta_1 - a) + v(\theta_2 - b)$$

Find θ_1 and θ_2 in the red circle
e **minimizing** $L(\theta)$

$$(\theta_1 - a)^2 + (\theta_2 - b)^2 \leq d^2$$



Third Derivation Method: Gradient Descent

Red Circle: (If the radius is small)


$$L(\theta) \approx \cancel{s} + u \frac{\theta_1 - a}{\Delta \theta_1} + v \frac{\theta_2 - b}{\Delta \theta_2}$$

Find θ_1 and θ_2 in the red circle
e **minimizing** $L(\theta)$

$$\frac{(\theta_1 - a)^2}{\Delta \theta_1^2} + \frac{(\theta_2 - b)^2}{\Delta \theta_2^2} \leq d^2$$

To minimize $L(\theta)$

$$\begin{bmatrix} \Delta \theta_1 \\ \Delta \theta_2 \end{bmatrix} = -\eta \begin{bmatrix} u \\ v \end{bmatrix} \quad \Rightarrow \quad \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} a \\ b \end{bmatrix} - \eta \begin{bmatrix} u \\ v \end{bmatrix}$$



Third Derivation Method: Gradient Descent

Based on Taylor Series:

If the red circle is ***small enough***, in the red circle

constant

$$L(\theta) \approx s + u(\theta_1 - a) + v(\theta_2 - b)$$

$$s = L(a, b)$$

$$u = \frac{\partial L(a, b)}{\partial \theta_1}, v = \frac{\partial L(a, b)}{\partial \theta_2}$$

Find θ_1 and θ_2 yielding the smallest value of $L(\theta)$ in the circle

$$\begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} a \\ b \end{bmatrix} - \eta \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} a \\ b \end{bmatrix} - \eta \begin{bmatrix} \frac{\partial L(a, b)}{\partial \theta_1} \\ \frac{\partial L(a, b)}{\partial \theta_2} \end{bmatrix}$$

This is gradient descent.

Not satisfied if the red circle (learning rate) is not small enough

You can consider the second order term, e.g. Newton's method.

Stochastic Gradient Descent

- Consider minimizing an average of functions

$$\min_x \frac{1}{m} \sum_{i=1}^m f_i(x)$$

- Gradient descent would repeat

$$x^{(k)} = x^{(k-1)} - t_k \cdot \frac{1}{m} \sum_{i=1}^m \nabla f_i(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

- In Stochastic Gradient Descent (SGD) (a.k.a. incremental gradient descent)

$$x^{(k)} = x^{(k-1)} - t_k \cdot \nabla f_{i_k}(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

where $i_k \in \{1, \dots, m\}$ is some chosen index at iteration k

Stochastic Gradient Descent

- **Two rules for choosing i_k at iteration k :**
 - Randomized rule: choose i_k from $\{1, \dots, m\}$ uniformly at random
 - Cyclic rule: choose $i_k=1, 2, \dots, m, 1, 2, \dots, m, \dots$

- **Main appeal of SGD**
 - Iteration cost is independent of m , the number of functions
 - Can also be a big savings in terms of memory usage

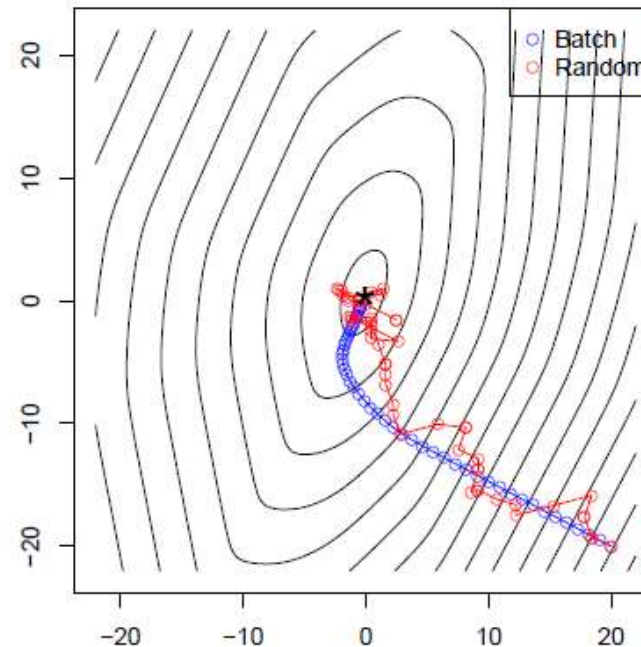
Stochastic Gradient Descent

Given $(x_i, y_i) \in \mathbb{R}^p \times \{0, 1\}$, $i = 1, \dots, n$, recall **logistic regression**:

$$\min_{\beta} \frac{1}{n} \sum_{i=1}^n \underbrace{\left(-y_i x_i^T \beta + \log(1 + \exp(x_i^T \beta)) \right)}_{f_i(\beta)}$$

Standard in SGD is to use **diminishing step sizes**, e.g., $t_k = 1/k$

Small example with $n = 10$, $p = 2$ to show the “classic picture” for batch versus stochastic methods:



Blue: batch steps, $O(np)$

Red: stochastic steps, $O(p)$

Rule of thumb for stochastic methods:

- generally thrive far from optimum
- generally struggle close to optimum

Mini-Batches Stochastic Gradient Descent

- We choose a random subset $I_k \subseteq \{1, \dots, m\}$, $|I_k| = b \ll m$, to repeat

$$x^{(k)} = x^{(k-1)} - t_k \cdot \frac{1}{b} \sum_{i \in I_k} \nabla f_i(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

- Example) Consider the problem

$$\min_{\beta} \frac{1}{n} \sum_{i=1}^n \left(-y_i x_i^T \beta + \log(1 + e^{x_i^T \beta}) \right) + \frac{\lambda}{2} \|\beta\|_2^2$$




Full gradient computation is $\nabla f(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - p_i(\beta)) x_i + \lambda \beta$.

Comparison between methods:

- One batch update costs $O(np)$
- One mini-batch update costs $O(bp)$
- One stochastic update costs $O(p)$

Comparison of Gradient Descent Methods

Example with $n = 10,000$, $p = 20$, all methods use fixed step sizes:



Read More About

- **Constraint optimization and KKT conditions**
- **Numerical optimization methods**
 - Genetic algorithm
 - Simulated annealing
 - Newton-Raphson method