



# Introduction to AI for postgraduate students

Lecture Note 8  
Optimization for Training Deep Models

**POSTECH**



# Agenda



- **How learning differs from pure optimization**
- Challenges in neural network optimization
- Basic algorithms
- Parameter initialization strategies
- Algorithms with adaptive learning rates

# How Learning Differs from Pure Optimization

- ML usually acts **indirectly**.
- In classical optimization, we optimize the **performance measure  $P$** .
- In ML,  $P$  is often **intractable**, and is defined with respect to **unobserved** test sets.
- In ML, we reduce a **different cost function  $J(\theta)$**  in the hope that doing so will improve  $P$ .
- The focus is on the **unregularized and supervised learning case**, where  $J$  is typically determined:

$$J(\theta) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{\text{data}}} L(f(\mathbf{x}; \theta), y),$$

- However, the ultimate goal is to minimize the **generalization error**, which is defined by:

$$J^*(\theta) = \mathbb{E}_{(\mathbf{x}, y) \sim p_{\text{data}}} L(f(\mathbf{x}; \theta), y).$$

# Empirical Risk Minimization

$$J(\theta) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{\text{data}}} L(f(\mathbf{x}; \theta), y),$$

Empirical  
risk

$$J^*(\theta) = \mathbb{E}_{(\mathbf{x}, y) \sim p_{\text{data}}} L(f(\mathbf{x}; \theta), y).$$

Risk

- The training process minimizing an **empirical risk** is very similar to straightforward optimization.
- But often, we use **different** approach in minimizing the empirical risk with ML.
  - **Early stopping, regularization**, etc.
  - Cf) In classical optimization, the optimization ends typically when the **gradient is very small**.
- It's because, the **ultimate goal** of ML is to minimize the **risk (generalization error)**, not the empirical risk.

# Surrogate Loss Function

- Sometimes the loss function we actually care about is **not** one that can be optimized efficiently.
  - Exactly minimizing 0-1 loss is typically **intractable**.
- **Surrogate loss function**
  - Acts as a **proxy** of the original goal but has **advantages**.
  - E.g., **negative log-likelihood** is used as a surrogate for the 0-1 loss.
  - We can improve the **robustness of the classifier by further** with a surrogate loss by pushing the classes apart from each other

# Batch and Minibatch Algorithms

- One major difference from the classical optimization in ML is that we generally use a **minibatch** algorithm.
  - The full batch provides a more accurate estimation of the gradient, but it is **not the best for ML** training.

- For instance, in maximum likelihood problems, our training goal is to minimize

$$J(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x}, y \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}, y; \boldsymbol{\theta})$$

who has the gradient as  $\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x}, y \sim \hat{p}_{\text{data}}} \nabla_{\boldsymbol{\theta}} \log p_{\text{model}}(\mathbf{x}, y; \boldsymbol{\theta})$

- We typically employ the **sample mean** theory to find the solution on  $\boldsymbol{\theta}$ :

$$\boldsymbol{\theta}_{\text{ML}} = \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta})$$

# Batch and Minibatch Algorithms

- Why we have to use a minibatch?
  - Full batch may be computationally **prohibited**.
  - Generalization error may be **suppressed by using minibatches**, because of the increased noise in the gradient estimation.
- Considerations needed to determine the batch size
  - Larger batches provide a **more accurate estimate** of the gradient, but with **less than linear** returns.
  - Multicore architectures are usually **underutilized** by **extremely small batches**.
  - If all examples in the batch are to be processed in parallel (as is typically the case), then the **amount of memory scales with the batch size**.
  - When using GPUs, it is common for **power of 2 batch sizes** to offer **better runtime**. Typical power of 2 batch sizes range from 32 to 256, with 16 sometimes being attempted for large models.
  - Small batches can offer a **regularizing effect**, perhaps due to the **noise** they add to the learning process.
  - Training with such a small batch size might require a **small learning rate** to maintain stability because of the high variance in the estimate of the gradient.
  - Minibatches should be selected **randomly**, and they should be **independent** as much as possible.

# Agenda



- How learning differs from pure optimization
- **Challenges in neural network optimization**
- Basic algorithms
- Parameter initialization strategies
- Algorithms with adaptive learning rates



# Challenge: III -Conditioning

- In Chapter 4 we learnt the followings:


- The **Taylor approximation** of  $f(\mathbf{x})$  around the **current** point  $\mathbf{x}^{(0)}$ :

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^\top \mathbf{g} + \frac{1}{2}(\mathbf{x} - \mathbf{x}^{(0)})^\top \mathbf{H}(\mathbf{x} - \mathbf{x}^{(0)})$$

- The **gradient descent update** gives us  $\mathbf{x} = \mathbf{x}^{(0)} - \epsilon \mathbf{g}$ , at which the cost function becomes

$$f(\mathbf{x}^{(0)} - \epsilon \mathbf{g}) \approx f(\mathbf{x}^{(0)}) - \epsilon \mathbf{g}^\top \mathbf{g} + \frac{1}{2} \epsilon^2 \mathbf{g}^\top \mathbf{H} \mathbf{g}$$

- We expect  $f(\mathbf{x}^{(0)}) > f(\mathbf{x}^{(0)} - \epsilon \mathbf{g})$ , so there will be a problem if  $\frac{1}{2} \epsilon^2 \mathbf{g}^\top \mathbf{H} \mathbf{g} > \epsilon \mathbf{g}^\top \mathbf{g}$ .
- In many cases,  $\mathbf{g}^\top \mathbf{H} \mathbf{g}$  grows easily.
  - Learning becomes very slow despite the presence of a strong gradient.

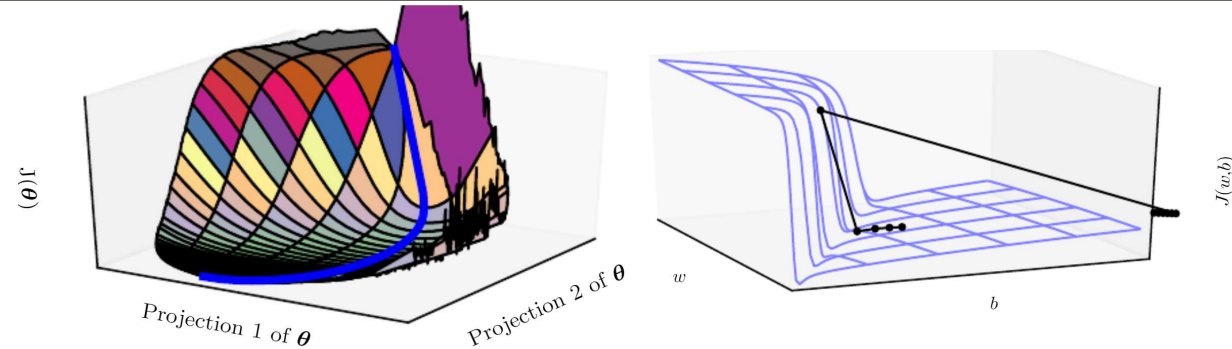


# Challenges: Others

- **Zero gradient** at non-optimal points
  - Local minima
  - Saddle points and other flat regions
  - **SGD** can be a solution
- Cliffs and Exploding Gradients
  - Gradient clipping can be a solution
- Vanishing and exploding gradient problem
  - Some neural networks (e.g., RNN) contain a path consisting of **repeatedly multiplying** by a matrix  $\mathbf{W}$ .
  - With an eigen decomposition notation  $\mathbf{W} = \mathbf{V}\text{diag}(\boldsymbol{\lambda})\mathbf{V}^{-1}$ , after  $t$  steps, the total multiplying factor becomes:


$$\mathbf{W}^t = (\mathbf{V}\text{diag}(\boldsymbol{\lambda})\mathbf{V}^{-1})^t = \mathbf{V}\text{diag}(\boldsymbol{\lambda})^t\mathbf{V}^{-1}$$

- If eigen values are smaller than 1 in magnitude → **vanishing** gradient problem
- If eigen values are larger than 1 in magnitude → **exploding** gradient problem → gradient clipping.



# Challenges: Others (Cont'd)

- Inexact gradients
  - Use a **surrogate loss function** that is easier to approximate than the true loss
- Poor initialization
  - In many cases, neural networks **do not** arrive at any critical points such as local minima, saddle points, etc.
  - But still, due to **poor initialization**, the learning result can be poor.
- Theoretical limits of optimization theory
  - e.g., discrete-valued units



# Agenda



- How learning differs from pure optimization
- Challenges in neural network optimization
- **Basic algorithms**
- Parameter initialization strategies
- Algorithms with adaptive learning rates

# Stochastic Gradient Descent

- SGD and its variants are the **most used** optimization algorithms for deep learning.
- It's possible to obtain an **unbiased** estimate of the gradient by taking the average gradient on a minibatch of  $m$  examples drawn **i.i.d.** from the data-generating distribution.

---

**Algorithm 8.1** Stochastic gradient descent (SGD) update

---

**Require:** Learning rate schedule  $\epsilon_1, \epsilon_2, \dots$

**Require:** Initial parameter  $\theta$

$k \leftarrow 1$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

    Apply update:  $\theta \leftarrow \theta - \epsilon_k \hat{\mathbf{g}}$

$k \leftarrow k + 1$

**end while**

---

# Stochastic Gradient Descent

- Scheduling the **learning rate** is very important.
  - We **gradually decrease** the learning rate over time.
  - This is because the SGD gradient estimator introduces a source of **noise** that does not vanish even when we arrive at a minimum.
- Sufficient condition to guarantee **convergence** of SGD:

$$\sum_{k=1}^{\infty} \epsilon_k = \infty, \quad \text{and}$$

$$\sum_{k=1}^{\infty} \epsilon_k^2 < \infty.$$

$\alpha = \frac{k}{\tau}$ , where  $\tau$  is set to the number of iterations required to make a few hundred passes through the training set.

- Typical choice of  $\epsilon_k$  is:

$$\epsilon_k = (1 - \alpha)\epsilon_0 + \alpha\epsilon_{\tau}$$

$\epsilon_{\tau}$  should be set to roughly **1 %** the value of  $\epsilon_0$ .

If it is **too large**, the learning curve will show **violent oscillations**.  
 If the learning rate is **too low**, learning proceeds **slowly**, and learning may **become stuck with a high cost value**.

# Momentum

- Drawback of the SGD: learning can be **slow**.
- **Momentum**
  - **Accelerates** learning, especially in the face of **high curvature**, **small but consistent** gradients, or **noisy** gradients.
  - Momentum in physics = mass × velocity.
  - Assuming unit mass, the momentum is the **velocity** here, which is defined by

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \nabla_{\boldsymbol{\theta}} \left( \frac{1}{m} \sum_{i=1}^m L(\mathbf{f}(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}) \right)$$

- $\alpha$ : hyperparameter
- Tries to maintain the previous velocity

- Incremental change in velocity

- Then, the parameter update is given by

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{v}$$

# SGD with Momentum

---

**Algorithm 8.2** Stochastic gradient descent (SGD) with momentum

---

**Require:** Learning rate  $\epsilon$ , momentum parameter  $\alpha$

**Require:** Initial parameter  $\theta$ , initial velocity  $v$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$ .

    Compute velocity update:  $\mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \mathbf{g}$ .

    Apply update:  $\theta \leftarrow \theta + \mathbf{v}$ .

**end while**

---




# SGD with Momentum

- SGD with momentum can **mitigate** abrupt change in gradients
  - **Fast and reliable** learning in the face of high curvature.
- SGD with momentum can also **accelerate** the update when the gradients are all lying in the same direction.
  - If the gradient's direction is always in the same direction, the accumulation of the velocity gives us the converged velocity  $\mathbf{v}$  where
$$\mathbf{v} = \alpha \mathbf{v} - \epsilon \mathbf{g} \rightarrow \mathbf{v} = -\frac{\epsilon \mathbf{g}}{1-\alpha}, \quad \mathbf{g} \left( = \nabla_{\theta} \left( \frac{1}{m} \sum_{i=1}^m L(\mathbf{f}(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)}) \right) \right)$$
  - So, the magnitude of  $\mathbf{v}$  becomes  $\frac{\epsilon \|\mathbf{g}\|}{1-\alpha}$ .
  - Typical choice is  $\alpha = \{0.5, 0.9, 0.99\}$ .
  - With  $\alpha = 0.9$ , we get  $\frac{1}{1-\alpha} = 10$ , so **10 times faster** update is possible compared to the original SGD.

# SGD with Momentum

The contour lines depict a quadratic loss function with a **poorly conditioned** Hessian matrix

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \nabla_{\theta} \left( \frac{1}{m} \sum_{i=1}^m L(\mathbf{f}(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)}) \right)$$



# Agenda



- How learning differs from pure optimization
- Challenges in neural network optimization
- Basic algorithms
- **Parameter initialization strategies**
- Algorithms with adaptive learning rates

# Parameter Initialization Strategies

- Deep learning training algorithms usually are iterative and thus require the user to specify some **initial point**.
- Training deep models is **strongly affected** by the choice of initializations.
  - Initial points determine whether the **algorithm converges at all**.
  - Initial points determine **how quickly** learning converges, and whether it converges to a point with **high or low** cost.
  - Initial points affect the **generalization error**.
- Modern initialization strategies **are simple and heuristic**.
  - Designing improved initialization strategies is a **difficult** task, because neural network optimization is not yet well understood.
  - Although some initial points are good in the optimization perspectives, but **detrimental** from the viewpoint of generalization error.

# Parameter Initialization Strategies

- Break symmetry
  - Initialize each unit to compute a **different** function from all the **other** units.
  - Gram-Schmidt orthogonalization can be used to find a symmetry-breaking initial weight matrix.
  - But random initialization works **well** even with **low** computational complexity.
- Choosing the magnitudes of initial points
  - **Large** initial weights are good because
    - ✓ We can expect a **stronger** symmetry-breaking effect
    - ✓ We can **avoid losing** signal during forward or back-propagation.
  - **Too large** initial weights are not good because
    - ✓ There can be **exploding gradients** problems.
    - ✓ They result in **saturation** after the activation function.
- Popular choice: sampling each unit from

$$W_{i,j} \sim U\left(-\sqrt{\frac{6}{m+n}}, \sqrt{\frac{6}{m+n}}\right) \quad \text{OR} \quad U\left(-\frac{1}{\sqrt{m}}, \frac{1}{\sqrt{m}}\right)$$

- where  $m$ : # of input units,  $n$ : # of output units

# Agenda



- How learning differs from pure optimization
- Challenges in neural network optimization
- Basic algorithms
- Parameter initialization strategies
- **Algorithms with adaptive learning rates**

# Algorithms with Adaptive Learning Rates

- **Learning rate** is very important to train deep models robustly and reliably.
- The cost is often **highly sensitive** to some directions in parameter space and insensitive to others.
  - Having the **same learning rate** for all parameters is not good.
- The momentum algorithm can mitigate these issues to some extent.
- Are there any **other** approaches?
  - Idea is to **separate learning rate for each parameter** and **automatically** adapt these learning rates throughout the course of learning.

# AdaGrad



---

## Algorithm 8.4 The AdaGrad algorithm

---

**Require:** Global learning rate  $\epsilon$

**Require:** Initial parameter  $\theta$

**Require:** Small constant  $\delta$ , perhaps  $10^{-7}$ , for numerical stability

Initialize gradient accumulation variable  $\mathbf{r} = \mathbf{0}$

**while** stopping criterion not met **do**

Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$ .

Accumulate squared gradient:  $\mathbf{r} \leftarrow \mathbf{r} + \mathbf{g} \odot \mathbf{g}$ . Accumulate the squares of the gradients

Compute update:  $\Delta \theta \leftarrow -\frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \mathbf{g}$ . (Division and square root applied element-wise)

Apply update:  $\theta \leftarrow \theta + \Delta \theta$ .

Larger gradient values result in faster decaying of the learning rate

**end while**

---



# RMSProp



---

## Algorithm 8.5 The RMSProp algorithm

---

**Require:** Global learning rate  $\epsilon$ , decay rate  $\rho$

**Require:** Initial parameter  $\theta$

**Require:** Small constant  $\delta$ , usually  $10^{-6}$ , used to stabilize division by small numbers

Initialize accumulation variables  $\mathbf{r} = 0$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$ .

    Accumulate squared gradient:  $\mathbf{r} \leftarrow \rho \mathbf{r} + (1 - \rho) \mathbf{g} \odot \mathbf{g}$ .

    Compute parameter update:  $\Delta \theta = -\frac{\epsilon}{\sqrt{\delta + \mathbf{r}}} \odot \mathbf{g}$ . ( $\frac{1}{\sqrt{\delta + \mathbf{r}}}$  applied element-wise)

**Moving average** instead of accumulation, which can avoid too fast decaying of the learning rate.

    Apply update:  $\theta \leftarrow \theta + \Delta \theta$ .

**end while**

---

# Adam



---

**Algorithm 8.7** The Adam algorithm

---

**Require:** Step size  $\epsilon$  (Suggested default: 0.001)

**Require:** Exponential decay rates for moment estimates,  $\rho_1$  and  $\rho_2$  in  $[0, 1)$ .  
(Suggested defaults: 0.9 and 0.999 respectively)

**Require:** Small constant  $\delta$  used for numerical stabilization (Suggested default:  $10^{-8}$ )

**Require:** Initial parameters  $\theta$

Initialize 1st and 2nd moment variables  $\mathbf{s} = \mathbf{0}$ ,  $\mathbf{r} = \mathbf{0}$

Initialize time step  $t = 0$

**while** stopping criterion not met **do**

Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

$t \leftarrow t + 1$

Update biased first moment estimate:  $\mathbf{s} \leftarrow \rho_1 \mathbf{s} + (1 - \rho_1) \mathbf{g}$

Update biased second moment estimate:  $\mathbf{r} \leftarrow \rho_2 \mathbf{r} + (1 - \rho_2) \mathbf{g} \odot \mathbf{g}$

Correct bias in first moment:  $\hat{\mathbf{s}} \leftarrow \frac{\mathbf{s}}{1 - \rho_1^t}$

Correct bias in second moment:  $\hat{\mathbf{r}} \leftarrow \frac{\mathbf{r}}{1 - \rho_2^t}$

Compute update:  $\Delta \theta = -\epsilon \frac{\hat{\mathbf{s}}}{\sqrt{\hat{\mathbf{r}} + \delta}}$  (operations applied element-wise)

Apply update:  $\theta \leftarrow \theta + \Delta \theta$

**end while**

---

Momentum method

Moving average of the squared gradients

# Choosing the Right Optimization Algorithm

- There is currently **no consensus** on "which algorithm should be the best for given problem and dataset".
- The most popular optimization algorithms actively in use include **SGD**, **SGD with momentum**, **RMSProp**, **RMSProp with momentum**, **AdaDelta**, and **Adam**.