


Numpy

- Numpy is the core library for scientific computing in Python.
- It provides a high-performance multidimensional array object, and tools for working with these arrays.

To use Numpy, we first need to import the `numpy` package:

```
 import numpy as np
```

Numpy: Array

```
▶ a = np.array([1, 2, 3]) # Create a rank 1 array
print(type(a), a.shape, a[0], a[1], a[2])
a[0] = 5 # Change an element of the array
print(a)
```

```
<class 'numpy.ndarray'> (3,) 1 2 3
[5 2 3]
```

```
[ ] b = np.array([[1,2,3],[4,5,6]]) # Create a rank 2 array
print(b)
```

```
[[1 2 3]
 [4 5 6]]
```

```
▶ print(b.shape)
print(b[0, 0], b[0, 1], b[1, 0])
```

```
▶ (2, 3)
1 2 4
```



Numpy: Array

```
[ ] a = np.zeros((2,2)) # Create an array of all zeros  
print(a)
```

```
[[0. 0.]  
 [0. 0.]]
```

```
▶ b = np.ones((1,2)) # Create an array of all ones  
print(b)
```

```
👤 [[1. 1.]]
```

```
[ ] c = np.full((2,2), 7) # Create a constant array  
print(c)
```

```
[[7 7]  
 [7 7]]
```

```
[ ] d = np.eye(2) # Create a 2x2 identity matrix  
print(d)
```

```
[[1. 0.]  
 [0. 1.]]
```

```
▶ e = np.random.random((2,2)) # Create an array filled with random values  
print(e)
```

```
👤 [[0.8690054 0.57244319]  
 [0.29647245 0.81464494]]
```

Numpy: Array indexing

- **Slicing:** Similar to Python lists, numpy arrays can be sliced.
- Since arrays may be multidimensional, you must specify a slice for each dimension of the array:

```
import numpy as np

# Create the following rank 2 array with shape (3, 4)
# [[ 1  2  3  4]
#  [ 5  6  7  8]
#  [ 9 10 11 12]]
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])

# Use slicing to pull out the subarray consisting of the first 2 rows
# and columns 1 and 2; b is the following array of shape (2, 2):
# [[2 3]
#  [6 7]]
b = a[:2, 1:3]
print(b)
```

```
[[2 3]
 [6 7]]
```

A slice of an array is a view into the same data, so modifying it will modify the original array.

```
[ ] print(a[0, 1])
    b[0, 0] = 77 # b[0, 0] is the same piece of data as a[0, 1]
    print(a[0, 1])
```

```
2
77
```



Numpy: Array indexing

- Two ways of accessing the data in the middle row of the array.
- Mixing integer indexing with slices yields an array of lower rank, while using only slices yields an array of the same rank as the original array:

```
[ ] # Create the following rank 2 array with shape (3, 4)
a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print(a)
```

```
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]
```

```
[ ] row_r1 = a[1, :]    # Rank 1 view of the second row of a
row_r2 = a[1:2, :]    # Rank 2 view of the second row of a
row_r3 = a[[1], :]    # Rank 2 view of the second row of a
print(row_r1, row_r1.shape)
print(row_r2, row_r2.shape)
print(row_r3, row_r3.shape)
```

```
[5 6 7 8] (4,)
[[5 6 7 8]] (1, 4)
[[5 6 7 8]] (1, 4)
```

```
[ ] # We can make the same distinction when accessing columns of an array:
col_r1 = a[:, 1]
col_r2 = a[:, 1:2]
print(col_r1, col_r1.shape)
print()
print(col_r2, col_r2.shape)
```

```
[ 2  6 10] (3,)
```

```
[[ 2]
 [ 6]
 [10]] (3, 1)
```



Numpy: Array indexing

```
▶ a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print(a)
print()

row_r1 = a[1, :] # Rank 1 view of the second row of a
print(row_r1, row_r1.shape)
row_r1[0]=77
print(row_r1)
print(a)
print()
```

```
↪ [[ 1  2  3  4]
   [ 5  6  7  8]
   [ 9 10 11 12]]
```

```
[5 6 7 8] (4,)
[77 6 7 8]
[[ 1  2  3  4]
 [77  6  7  8]
 [ 9 10 11 12]]
```

```
[19] a = np.array([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print(a)
print()

row_r2 = np.array(a[1, :]) # Rank 2 view of the second row of a
print(row_r2, row_r2.shape)
row_r2[0] = 99
print(row_r2)
print(a)
print()
```

```
↪ [[ 1  2  3  4]
   [ 5  6  7  8]
   [ 9 10 11 12]]
```

```
[5 6 7 8] (4,)
[99 6 7 8]
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]
```



Numpy: Array indexing

- One useful trick with integer array indexing is selecting or mutating one element from each row of a matrix:

```
# Create a new array from which we will select elements
a = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
print(a)
```

```
[[ 1  2  3]
 [ 4  5  6]
 [ 7  8  9]
 [10 11 12]]
```

```
[ ] # Create an array of indices
b = np.array([0, 2, 0, 1])

# Select one element from each row of a using the indices in b
print(a[np.arange(4), b]) # Prints "[ 1  6  7 11]"
```

[0,1,2,3]

```
[ 1  6  7 11]
```

```
[ ] # Mutate one element from each row of a using the indices in b
a[np.arange(4), b] += 10
print(a)
```

```
[[11  2  3]
 [ 4  5 16]
 [17  8  9]
 [10 21 12]]
```



Numpy: Array indexing

- **Boolean array indexing:** Boolean array indexing lets you pick out arbitrary elements of an array.

```
▶ import numpy as np

a = np.array([[1,2], [3, 4], [5, 6]])

bool_idx = (a > 2) # Find the elements of a that are bigger than 2;
                  # this returns a numpy array of Booleans of the same
                  # shape as a, where each slot of bool_idx tells
                  # whether that element of a is > 2.

print(bool_idx)
```

```
[[False False]
 [ True  True]
 [ True  True]]
```

```
[ ] # We use boolean array indexing to construct a rank 1 array
    # consisting of the elements of a corresponding to the True values
    # of bool_idx
    print(a[bool_idx])

    # We can do all of the above in a single concise statement:
    print(a[a > 2])
```

```
[3 4 5 6]
[3 4 5 6]
```



Numpy: Datatypes

- Numpy tries to guess a datatype when you create an array, but functions that construct arrays usually also include an optional argument to explicitly specify the datatype.

```
▶ x = np.array([1, 2]) # Let numpy choose the datatype  
y = np.array([1.0, 2.0]) # Let numpy choose the datatype  
z = np.array([1, 2], dtype=np.int64) # Force a particular datatype  
  
print(x.dtype, y.dtype, z.dtype)
```

```
⊙ int64 float64 int64
```

Numpy: Array math

```
▶ x = np.array([[1,2],[3,4]], dtype=np.float64)
  y = np.array([[5,6],[7,8]], dtype=np.float64)

  # Elementwise sum; both produce the array
  print(x + y)
  print(np.add(x, y))
```

```
ⓘ [[ 6.  8.]
   [10. 12.]]
   [[ 6.  8.]
   [10. 12.]]
```

```
[ ] # Elementwise difference; both produce the array
    print(x - y)
    print(np.subtract(x, y))
```

```
[[ -4. -4.]
 [ -4. -4.]]
[[ -4. -4.]
 [ -4. -4.]]
```

```
[ ] # Elementwise product; both produce the array
    print(x * y)
    print(np.multiply(x, y))
```

```
[[ 5. 12.]
 [21. 32.]]
[[ 5. 12.]
 [21. 32.]]
```

```
[ ] # Elementwise division; both produce the array
    # [[ 0.2      0.33333333]
    # [ 0.42857143  0.5      ]]
    print(x / y)
    print(np.divide(x, y))
```

```
[[0.2      0.33333333]
 [0.42857143 0.5      ]]
[[0.2      0.33333333]
 [0.42857143 0.5      ]]
```

```
[ ] # Elementwise square root; produces the array
    # [[ 1.      1.41421356]
    # [ 1.73205081  2.      ]]
    print(np.sqrt(x))
```

```
[[1.      1.41421356]
 [1.73205081 2.      ]]
```

Numpy: Array math

- We use the dot function to compute inner products of vectors, to multiply a vector by a matrix, and to multiply matrices.
- dot is available both as a function in the numpy module and as an instance method of array objects:

```
v = np.array([9, 10])
w = np.array([11, 12])

# Inner product of vectors; both produce 219
print(v.dot(w))
print(np.dot(v, w))
```

```
219
219
```

Numpy: Array math

- You can also use the `@` operator which is equivalent to numpy's `dot` operator.

```
x = np.array([[1,2],[3,4]])  
y = np.array([[5,6],[7,8]])
```

```
v = np.array([9,10])  
w = np.array([11, 12])
```

```
[ ] print(v @ w)
```

```
219
```

```
▶ # Matrix / vector product; both produce the rank 1 array [29 67]  
print(x.dot(v))  
print(np.dot(x, v))  
print(x @ v)
```

```
⊞ [29 67]  
[29 67]  
[29 67]
```

```
[ ] # Matrix / matrix product; both produce the rank 2 array  
# [[19 22]  
#  [43 50]]  
print(x.dot(y))  
print(np.dot(x, y))  
print(x @ y)
```

```
[[19 22]  
 [43 50]]  
[[19 22]  
 [43 50]]  
[[19 22]  
 [43 50]]
```

Numpy: Array math

- Numpy provides many useful functions for performing computations on arrays; one of the most useful is sum:

```
▶ x = np.array([[1,2],[3,4]])  
  
print(np.sum(x)) # Compute sum of all elements; prints "10"  
print(np.sum(x, axis=0)) # Compute sum of each column; prints "[4 6]"  
print(np.sum(x, axis=1)) # Compute sum of each row; prints "[3 7]"
```

```
10  
[4 6]  
[3 7]
```

Numpy: Array math

▪ Transpose operation

```
▶ print(x)  
print("transpose#\n", x.T)
```

```
⦿ [[1 2]  
   [3 4]]  
transpose  
[[1 3]  
 [2 4]]
```

```
[ ] v = np.array([[1,2,3]])  
print(v )  
print("transpose#\n", v.T)
```

```
[[1 2 3]]  
transpose  
[[1]  
 [2]  
 [3]]
```

Numpy: Broadcasting

- **Broadcasting is a powerful mechanism that allows numpy to work with arrays of different shapes when performing arithmetic operations.**

```
# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = np.empty_like(x) # Create an empty matrix with the same shape as x

# Add the vector v to each row of the matrix x with an explicit loop
for i in range(4):
    y[i, :] = x[i, :] + v

print(y)
```

```
[[ 2  2  4]
 [ 5  5  7]
 [ 8  8 10]
 [11 11 13]]
```

```
vv = np.tile(v, (4, 1)) # Stack 4 copies of v on top of each other
print(vv)               # Prints "[[1 0 1]
                        #         [1 0 1]
                        #         [1 0 1]
                        #         [1 0 1]]"
```

```
[[1 0 1]
 [1 0 1]
 [1 0 1]
 [1 0 1]]
```

```
[ ] y = x + vv # Add x and vv elementwise
print(y)
```

```
[[ 2  2  4]
 [ 5  5  7]
 [ 8  8 10]
 [11 11 13]]
```

Numpy: Broadcasting

```
import numpy as np

# We will add the vector v to each row of the matrix x,
# storing the result in the matrix y
x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = np.array([1, 0, 1])
y = x + v # Add v to each row of x using broadcasting
print(y)
```

```
[[ 2  2  4]
 [ 5  5  7]
 [ 8  8 10]
 [11 11 13]]
```

- The line `y = x + v` works even though `x` has shape (4, 3) and `v` has shape (3,) due to broadcasting; this line works as if `v` actually had shape (4, 3), where each row was a copy of `v`, and the sum was performed elementwise.
- Broadcasting two arrays together follows these rules:
 - If the arrays do not have the same rank, prepend the shape of the lower rank array with 1s until both shapes have the same length.
 - The two arrays are said to be compatible in a dimension if they have the same size in the dimension, or if one of the arrays has size 1 in that dimension.
 - The arrays can be broadcast together if they are compatible in all dimensions.
 - After broadcasting, each array behaves as if it had shape equal to the elementwise maximum of shapes of the two input arrays.
 - In any dimension where one array had size 1 and the other array had size greater than 1, the first array behaves as if it were copied along that dimension


Numpy: Broadcasting

Here are some applications of broadcasting:

▶ # Compute outer product of vectors
`v = np.array([1,2,3])` # v has shape (3,)
`w = np.array([4,5])` # w has shape (2,)
 # To compute an outer product, we first reshape v to be a column
 # vector of shape (3, 1); we can then broadcast it against w to yield
 # an output of shape (3, 2), which is the outer product of v and w:


```
print(np.reshape(v, (3, 1)) * w)
```

▶ `[[4 5]
 [8 10]
 [12 15]]`



▶ # Add a vector to each row of a matrix
`x = np.array([[1,2,3], [4,5,6]])`
 # x has shape (2, 3) and v has shape (3,) so they broadcast to (2, 3),
 # giving the following matrix:

```
print(x + v)
```




▶ `[[2 4 6]
 [5 7 9]]`

Numpy: Broadcasting



```
# Add a vector to each column of a matrix
# x has shape (2, 3) and w has shape (2,).
# If we transpose x then it has shape (3, 2) and can be broadcast
# against w to yield a result of shape (3, 2); transposing this result
# yields the final result of shape (2, 3) which is the matrix x with
# the vector w added to each column. Gives the following matrix:
```


```
print((x.T + w).T)
```



```
[[ 5  6  7]
 [ 9 10 11]]
```


```
[ ] # Another solution is to reshape w to be a row vector of shape (2, 1);
# we can then broadcast it directly against x to produce the same
# output.
print(x + np.reshape(w, (2, 1)))
```

```
[[ 5  6  7]
 [ 9 10 11]]
```



Numpy: Broadcasting

▶ # Multiply a matrix by a constant:
x has shape (2, 3). Numpy treats scalars as arrays of shape ();
these can be broadcast together to shape (2, 3), producing the
following array:
print(x * 2)



[[2 4 6]
 [8 10 12]]


Matplotlib

- Matplotlib is a plotting library

```
[ ] import matplotlib.pyplot as plt
```

```
▶ # Compute the x and y coordinates for points on a sine curve  
x = np.arange(0, 3 * np.pi, 0.1)  
y = np.sin(x)  
  
# Plot the points using matplotlib  
plt.plot(x, y)
```

[<matplotlib.lines.Line2D at 0x7f0f3a0b4208>]




Matplotlib



```
y_sin = np.sin(x)  
y_cos = np.cos(x)
```

```
# Plot the points using matplotlib  
plt.plot(x, y_sin)  
plt.plot(x, y_cos)  
plt.xlabel('x axis label')  
plt.ylabel('y axis label')  
plt.title('Sine and Cosine')  
plt.legend(['Sine', 'Cosine'])
```

<matplotlib.legend.Legend at 0x7f0f39c04780>



Matplotlib



```
) # Compute the x and y coordinates for points on sine and cosine curves
x = np.arange(0, 3 * np.pi, 0.1)
y_sin = np.sin(x)
y_cos = np.cos(x)

# Set up a subplot grid that has height 2 and width 1,
# and set the first such subplot as active.
plt.subplot(2, 1, 1)

# Make the first plot
plt.plot(x, y_sin)
plt.title('Sine')

# Set the second subplot as active, and make the second plot.
plt.subplot(2, 1, 2)
plt.plot(x, y_cos)
plt.title('Cosine')

# Show the figure.
plt.show()
```

